

# Toward Bayesian Deep Learning

MI2RL 장령우

# Yarin Gal

- Godfather of modern Bayesian Deep Learning
- Ph.D thesis : “Uncertainty of Deep Learning”(<http://mlg.eng.cam.ac.uk/yarin/thesis/thesis.pdf>)

# Uncertainty in Deep Learning

- Bayes' Rule

$$f(x|y) = \frac{f(y|x)f(x)}{f(y)} \propto \mathcal{L}(y|x)f(x)$$

- $f(x|y)$  : posterior probability
- $f(x)$  : prior probability of  $X$
- $f(y)$  : prior probability of  $Y$  (normalizing constant)
- $f(y|x)$  : conditional probability given  $X$
- $\mathcal{L}(y|x)$  : likelihood

# Uncertainty in Deep Learning

- Types of uncertainty
  - **Noisy data** : Observations can be noisy d/t measurement imprecision -> **Aleatoric uncertainty**
  - **Uncertainty in model parameters** : Large number of models can be used to explain train data.  
Which model to use? -> **Epistemic uncertainty 1**
  - **Structural uncertainty** : What model structure to use? -> **Epistemic uncertainty 2**
- Predictive uncertainty = **Aleatoric uncertainty** + **Epistemic uncertainty**

(confidence we have in prediction)

- Epistemic : Came from “episteme”, Greek for “knowledge” -> “reducible unc-”
- Aleatoric : Came from “aleator”, Latin for “dice player” -> “irreducible unc-”

# Uncertainty in Deep Learning

- Deep learning model does not offer uncertainties!
- Does softmax output can be interpreted as model confidence? (classification)
  - A model can be wrong although having high softmax

# Uncertainty in Deep Learning

- Given a dataset  $X, Y$ , we look for the *posterior* distribution over the space of parameters by invoking Bayes' theorem:

$$p(\omega|X, Y) = \frac{p(Y|X, \omega)p(\omega)}{p(Y|X)}$$

- After training, we can predict new data as follows (inference):

$$p(y^*|x^*, X, Y) = \int p(y^*|x^*, \omega)p(\omega|X, Y)d\omega$$

- Model evidence:

$$p(Y|X) = \int p(Y|X, \omega)p(\omega)d\omega$$

# Uncertainty in Deep Learning

- True posterior  $p(\omega|X, Y)$  cannot be computed analytically.
- Instead, we define a *variational* distribution  $q_\theta(\omega)$  which is easy to evaluate.
- Thus, we minimize Kullback-Leibler (KL) divergence w.r.t  $\theta$  :

$$KL(q_\theta(\omega)||p(\omega|X, Y)) = \int q_\theta(\omega) \log \frac{q_\theta(\omega)}{p(\omega|X, Y)} d\omega$$

- Minimising KL divergence allows us for :

$$\begin{aligned} p(y^*|x^*, X, Y) &= \int p(y^*|x^*, \omega) p(\omega|X, Y) d\omega \\ &\approx \int p(y^*|x^*, \omega) q_\theta^*(\omega) d\omega =: q_\theta^*(y^*|x^*) \end{aligned}$$

# Uncertainty in Deep Learning

- KL divergence minimising == Maximising *evidence lower bound (ELBO)* w.r.t variational parameters defining  $q_{\theta}(\omega)$

$$\begin{aligned}\mathcal{L}_{VI}(\theta) &:= \int q_{\theta}(\omega) \log p(Y|X, \omega) \\ &\leq \log p(Y|X)\end{aligned}$$

where  $\log p(Y|X)$  is “log evidence”, and VI is “Variational Inference”



# Uncertainty in Deep Learning

- As

$$\begin{aligned} KL(q_{\theta}(\omega) || p(\omega|X, Y)) &\propto - \int q_{\theta}(\omega) \log p(Y|X, \omega) d\omega + KL(q_{\theta}(\omega) || p(\omega)) \\ &= - \sum_{i=1}^N \int q_{\theta}(\omega) \log p(y_i | f^{\omega}(x_i)) d\omega + KL(q_{\theta}(\omega) || p(\omega)) \end{aligned}$$

- We set

$$\mathcal{L}_{VI}(\theta) := KL(q_{\theta}(\omega) || p(\omega|X, Y))$$

# Uncertainty in Deep Learning

- Therefore,

$$\mathcal{L}_{VI}(\theta) := - \sum_{i=1}^N \int q_{\theta}(\omega) \log p(y_i | f^{\omega}(x_i)) d\omega + KL(q_{\theta}(\omega) || p(\omega))$$

- This has two problems:
  - First term on RHS is not tractable with more than single hidden layer.
  - This objective requires computation over the whole dataset, which is computationally expensive.
- Therefore, mini-batch optimisation (data sub-sampling)

$$\hat{\mathcal{L}}_{VI}(\theta) := - \frac{N}{M} \sum_{i \in S} \int q_{\theta}(\omega) \log p(y_i | f^{\omega}(x_i)) d\omega + KL(q_{\theta}(\omega) || p(\omega))$$

where S is random index set with size M

# Uncertainty in Deep Learning

- Data sub-sampling approximation forms an unbiased estimator, meaning

$$\mathbb{E}_S[\hat{\mathcal{L}}_{VI}(\theta)] = \mathcal{L}_{VI}(\theta)$$

- Monte Carlo is often used for variational inference. Yet there is gradient descent method:

---

**Algorithm 1** Minimise divergence between  $q_\theta(\omega)$  and  $p(\omega|X, Y)$

---

- 1: Given dataset  $\mathbf{X}, \mathbf{Y}$ ,
- 2: Define learning rate schedule  $\eta$ ,
- 3: Initialise parameters  $\theta$  randomly.
- 4: **repeat**
- 5:   Sample  $M$  random variables  $\hat{\epsilon}_i \sim p(\epsilon)$ ,  $S$  a random subset of  $\{1, \dots, N\}$  of size  $M$ .
- 6:   Calculate stochastic derivative estimator w.r.t.  $\theta$ :

$$\widehat{\Delta\theta} \leftarrow -\frac{N}{M} \sum_{i \in S} \frac{\partial}{\partial \theta} \log p(\mathbf{y}_i | \mathbf{f}^{g(\theta, \hat{\epsilon}_i)}(\mathbf{x}_i)) + \frac{\partial}{\partial \theta} \text{KL}(q_\theta(\omega) || p(\omega)).$$

- 7:   Update  $\theta$ :  
       $\theta \leftarrow \theta + \eta \widehat{\Delta\theta}$ .
  - 8: **until**  $\theta$  has converged.
-

# Uncertainty in Deep Learning

- Stochastic Regulation Technique (SRT) : Dropout !
- $\hat{\epsilon}_i$  : vector with dimension Q (input dimensionality) whose elements are random probabilities. (Dropout rate)
- $\hat{x} = x \odot \hat{\epsilon}_i$  where x is input vector.
- $\hat{h} = \sigma(\hat{x}M_1 + b) \odot \hat{\epsilon}_2$
- Output :  $\hat{y} = \hat{h}M_2$

# Uncertainty in Deep Learning

- Therefore,
$$\begin{aligned}\hat{\mathbf{y}} &= \hat{\mathbf{h}}\mathbf{M}_2 \\ &= (\mathbf{h} \odot \hat{\boldsymbol{\epsilon}}_2)\mathbf{M}_2 \\ &= (\mathbf{h} \cdot \text{diag}(\hat{\boldsymbol{\epsilon}}_2))\mathbf{M}_2 \\ &= \mathbf{h}(\text{diag}(\hat{\boldsymbol{\epsilon}}_2)\mathbf{M}_2) \\ &= \sigma(\hat{\mathbf{x}}\mathbf{M}_1 + \mathbf{b})(\text{diag}(\hat{\boldsymbol{\epsilon}}_2)\mathbf{M}_2) \\ &= \sigma((\mathbf{x} \odot \hat{\boldsymbol{\epsilon}}_1)\mathbf{M}_1 + \mathbf{b})(\text{diag}(\hat{\boldsymbol{\epsilon}}_2)\mathbf{M}_2) \\ &= \sigma(\mathbf{x}(\text{diag}(\hat{\boldsymbol{\epsilon}}_1)\mathbf{M}_1) + \mathbf{b})(\text{diag}(\hat{\boldsymbol{\epsilon}}_2)\mathbf{M}_2)\end{aligned}$$
- Which implies,
$$\frac{\partial}{\partial \theta} \hat{\mathcal{L}}_{\text{dropout}}(\theta) = \frac{1}{N_\tau} \frac{\partial}{\partial \theta} \hat{\mathcal{L}}_{\text{MC}}(\theta)$$
- Dropout as a Bayesian Approximation: Representing Model Uncertainty in Deep Learning(<https://arxiv.org/pdf/1506.02142.pdf>)

# Uncertainty in Deep Learning

- Dropout as a Bayesian Approximation: Representing Model Uncertainty in Deep Learning(<https://arxiv.org/pdf/1506.02142.pdf>)

- Dropout NN is equivalent as probabilistic Gaussian Process
- Uncertainty of dropout NN model is expressed as :

$$\log p(y^*|x^*, X, Y) \approx \log \left( \sum_{i=1}^T \exp(-\frac{1}{2}\tau ||y - \hat{y}_t||^2) \right) - \log T$$
$$- \frac{1}{2} \log 2\pi - \frac{1}{2} \log \tau^{-1}$$

- Gaussian process : A stochastic process that is composed with normal distributions:  $X_{(t_1, t_2, \dots, t_n)} = (X_{t_1}, X_{t_2}, \dots, X_{t_n})$

# What Uncertainties Do We Need in Bayesian Deep Learning for Computer Vision?

- Epistemic Uncertainty in Bayesian Deep Learning :

$$p(y = c|x, X, Y) \approx \frac{1}{T} \sum_{t=1}^T \text{Softmax}(f^{\widehat{W}_t}(x))$$

where uncertainty of this probability vector  $p$  is expressed as entropy  $H$ :

$$H(p) = - \sum_{c=1}^C p_c \log p_c$$

- Aleatoric Uncertainty :

$$\mathcal{L}_{NN}(\theta) = \frac{1}{N} \sum_{i=1}^N \frac{1}{2\sigma(x_i)^2} \|y_i - f(x_i)\|^2 + \frac{1}{2} \log \sigma(x_i)^2$$

# What Uncertainties Do We Need in Bayesian Deep Learning for Computer Vision?

- Combining aleatoric and epistemic uncertainty in one model :

$$[\hat{y}, \hat{\sigma}^2] = f^{\widehat{W}}(x)$$

- To summarize, with  $\hat{y}_t, \hat{\sigma}_t^2 = f^{\widehat{W}_t}(x)$ , with T models and

$$\text{Var}(y) \approx \frac{1}{T} \sum_{i=1}^T \hat{y}_t^2 - \left( \frac{1}{T} \sum_{t=1}^T \hat{y}_t \right)^2 + \frac{1}{T} \sum_{t=1}^T \hat{\sigma}_t^2$$

- Objective : To minimise

$$\mathcal{L}_{BNN}(\theta) = \frac{1}{D} \sum_i \frac{1}{2} \hat{\sigma}_i^{-2} ||y_i - \hat{y}_i||^2 + \frac{1}{2} \log \hat{\sigma}_i^2$$