

Toward Bayesian Deep Learning

MI2RL 장령우

Yarin Gal

- Rising star of modern Bayesian Deep Learning
- Ph.D thesis : “Uncertainty of Deep Learning”(<http://mlg.eng.cam.ac.uk/yarin/thesis/thesis.pdf>)

Uncertainty in Deep Learning

- Bayes' Rule

$$f(x|y) = \frac{f(y|x)f(x)}{f(y)} \propto \mathcal{L}(y|x)f(x)$$

- $f(x|y)$: **posterior** probability
- $f(x)$: **prior** probability of X
- $f(y)$: **evidence** (normalizing constant)
- $f(y|x)$: conditional probability given X (**likelihood**)
- $\mathcal{L}(y|x)$: **likelihood**

Uncertainty in Deep Learning

- Types of uncertainty
 - **Noisy data** : Observations can be noisy d/t measurement imprecision -> **Aleatoric uncertainty**
 - **Uncertainty in model parameters** : Large number of models can be used to explain train data.
Which model to use? -> **Epistemic uncertainty 1**
 - **Structural uncertainty** : What model structure to use? -> **Epistemic uncertainty 2**
- Predictive uncertainty = **Aleatoric uncertainty** + **Epistemic uncertainty**

(confidence we have in prediction)

- Epistemic : Came from “episteme”, Greek for “knowledge” -> “reducible unc-”
- Aleatoric : Came from “aleator”, Latin for “dice player” -> “irreducible unc-”

Uncertainty in Deep Learning

- Deep learning model does not offer uncertainties!
- Does softmax output can be interpreted as model confidence? (classification)
 - A model can be wrong although having high softmax

Uncertainty in Deep Learning

- Given a dataset X, Y , we look for the *posterior* distribution over the space of parameters by invoking Bayes' theorem:

$$p(\omega|X, Y) = \frac{p(Y|X, \omega)p(\omega)}{p(Y|X)}$$

- After training, we can predict new data as follows (inference):

$$p(y^*|x^*, X, Y) = \int p(y^*|x^*, \omega)p(\omega|X, Y)d\omega$$

- Model evidence:

$$p(Y|X) = \int p(Y|X, \omega)p(\omega)d\omega$$

Uncertainty in Deep Learning

- True posterior $p(\omega|X, Y)$ cannot be computed analytically.
- Instead, we define a *variational* distribution $q_\theta(\omega)$ which is easy to evaluate.
- Thus, we minimize Kullback-Leibler (KL) divergence w.r.t θ :

$$KL(q_\theta(\omega)||p(\omega|X, Y)) = \int q_\theta(\omega) \log \frac{q_\theta(\omega)}{p(\omega|X, Y)} d\omega$$

- Minimising KL divergence allows us for :

$$\begin{aligned} p(y^*|x^*, X, Y) &= \int p(y^*|x^*, \omega) p(\omega|X, Y) d\omega \\ &\approx \int p(y^*|x^*, \omega) q_\theta^*(\omega) d\omega =: q_\theta^*(y^*|x^*) \end{aligned}$$

Uncertainty in Deep Learning

- KL divergence minimising == Maximising *evidence lower bound (ELBO)* w.r.t variational parameters defining $q_{\theta}(\omega)$

$$\begin{aligned}\mathcal{L}_{VI}(\theta) &= \int q_{\theta}(\omega) \log p(Y|X, \omega) d\omega - KL\left(q_{\theta}(\omega) || p(\omega)\right) \\ &\leq \log p(Y|X)\end{aligned}$$

where $\log p(Y|X)$ is “log (model) evidence”, and VI is “variational inference”

- Maximising the first term (*expected log likelihood*) encourages model to explain data well, while minimising the second term (*prior KL*) encourages model to be as close as possible to the prior

Uncertainty in Deep Learning

- As

$$\begin{aligned} KL(q_{\theta}(\omega) || p(\omega | X, Y)) &\propto - \int q_{\theta}(\omega) \log p(Y | X, \omega) d\omega + KL(q_{\theta}(\omega) || p(\omega)) \\ &= - \sum_{i=1}^N \int q_{\theta}(\omega) \log p(y_i | f^{\omega}(x_i)) d\omega + KL(q_{\theta}(\omega) || p(\omega)) \end{aligned}$$

Uncertainty in Deep Learning

- Variational Inference :

$$\mathcal{L}_{VI}(\theta) := - \sum_{i=1}^N \int q_{\theta}(\omega) \log p(y_i | f^{\omega}(x_i)) d\omega + KL(q_{\theta}(\omega) || p(\omega))$$

- This has two problems:
 - First term on RHS is not tractable with more than single hidden layer.
 - This objective requires computation over the whole dataset, which is computationally expensive.
- Therefore, mini-batch optimisation (data sub-sampling)

$$\hat{\mathcal{L}}_{VI}(\theta) := - \frac{N}{M} \sum_{i \in S} \int q_{\theta}(\omega) \log p(y_i | f^{\omega}(x_i)) d\omega + KL(q_{\theta}(\omega) || p(\omega))$$

where S is random index set with size M

Uncertainty in Deep Learning

- Data sub-sampling approximation forms an unbiased estimator, meaning

$$\mathbb{E}_S[\hat{\mathcal{L}}_{VI}(\theta)] = \mathcal{L}_{VI}(\theta)$$

- Monte Carlo is often used for variational inference. Yet there is gradient descent method:

Algorithm 1 Minimise divergence between $q_\theta(\omega)$ and $p(\omega|X, Y)$

- 1: Given dataset \mathbf{X}, \mathbf{Y} ,
- 2: Define learning rate schedule η ,
- 3: Initialise parameters θ randomly.
- 4: **repeat**
- 5: Sample M random variables $\hat{\epsilon}_i \sim p(\epsilon)$, S a random subset of $\{1, \dots, N\}$ of size M .
- 6: Calculate stochastic derivative estimator w.r.t. θ :

$$\widehat{\Delta\theta} \leftarrow -\frac{N}{M} \sum_{i \in S} \frac{\partial}{\partial \theta} \log p(\mathbf{y}_i | \mathbf{f}^{g(\theta, \hat{\epsilon}_i)}(\mathbf{x}_i)) + \frac{\partial}{\partial \theta} \text{KL}(q_\theta(\omega) || p(\omega)).$$

- 7: Update θ :
 $\theta \leftarrow \theta + \eta \widehat{\Delta\theta}$.
 - 8: **until** θ has converged.
-

Uncertainty in Deep Learning

- Stochastic Regulation Technique (SRT) : Dropout !
- $\hat{\epsilon}_i$: vector with dimension Q (input dimensionality) whose elements are random probabilities. (Dropout rate)
- $\hat{x} = x \odot \hat{\epsilon}_i$ where x is input vector.
- $\hat{h} = \sigma(\hat{x}M_1 + b) \odot \hat{\epsilon}_2$
- Output : $\hat{y} = \hat{h}M_2$

Uncertainty in Deep Learning

- Therefore,
$$\begin{aligned}\hat{\mathbf{y}} &= \hat{\mathbf{h}}\mathbf{M}_2 \\ &= (\mathbf{h} \odot \hat{\boldsymbol{\epsilon}}_2)\mathbf{M}_2 \\ &= (\mathbf{h} \cdot \text{diag}(\hat{\boldsymbol{\epsilon}}_2))\mathbf{M}_2 \\ &= \mathbf{h}(\text{diag}(\hat{\boldsymbol{\epsilon}}_2)\mathbf{M}_2) \\ &= \sigma(\hat{\mathbf{x}}\mathbf{M}_1 + \mathbf{b})(\text{diag}(\hat{\boldsymbol{\epsilon}}_2)\mathbf{M}_2) \\ &= \sigma((\mathbf{x} \odot \hat{\boldsymbol{\epsilon}}_1)\mathbf{M}_1 + \mathbf{b})(\text{diag}(\hat{\boldsymbol{\epsilon}}_2)\mathbf{M}_2) \\ &= \sigma(\mathbf{x}(\text{diag}(\hat{\boldsymbol{\epsilon}}_1)\mathbf{M}_1) + \mathbf{b})(\text{diag}(\hat{\boldsymbol{\epsilon}}_2)\mathbf{M}_2)\end{aligned}$$
- Which implies,
$$\frac{\partial}{\partial \theta} \hat{\mathcal{L}}_{\text{dropout}}(\theta) = \frac{1}{N_\tau} \frac{\partial}{\partial \theta} \hat{\mathcal{L}}_{\text{MC}}(\theta)$$
- Dropout as a Bayesian Approximation: Representing Model Uncertainty in Deep Learning(<https://arxiv.org/pdf/1506.02142.pdf>)

Uncertainty in Deep Learning

- Dropout as a Bayesian Approximation: Representing Model Uncertainty in Deep Learning(<https://arxiv.org/pdf/1506.02142.pdf>)

- Dropout NN is equivalent as probabilistic Gaussian Process
- Uncertainty of dropout NN model is expressed as :

$$\log p(y^*|x^*, X, Y) \approx \log \left(\sum_{i=1}^T \exp(-\frac{1}{2}\tau ||y - \hat{y}_t||^2) \right) - \log T$$
$$- \frac{1}{2} \log 2\pi - \frac{1}{2} \log \tau^{-1}$$

- Gaussian process : A stochastic process that is composed with normal distributions: $X_{(t_1, t_2, \dots, t_n)} = (X_{t_1}, X_{t_2}, \dots, X_{t_n})$

What Uncertainties Do We Need in Bayesian Deep Learning for Computer Vision?

- Epistemic Uncertainty in Bayesian Deep Learning :

$$p(y = c|x, X, Y) \approx \frac{1}{T} \sum_{t=1}^T \text{Softmax}(f^{\widehat{W}_t}(x))$$

where uncertainty of this probability vector p is expressed as entropy H :

$$H(p) = - \sum_{c=1}^C p_c \log p_c$$

- Aleatoric Uncertainty :

$$\mathcal{L}_{NN}(\theta) = \frac{1}{N} \sum_{i=1}^N \frac{1}{2\sigma(x_i)^2} \|y_i - f(x_i)\|^2 + \frac{1}{2} \log \sigma(x_i)^2$$

What Uncertainties Do We Need in Bayesian Deep Learning for Computer Vision?

- Combining aleatoric and epistemic uncertainty in one model :

$$[\hat{y}, \hat{\sigma}^2] = f^{\widehat{W}}(x)$$

- To summarize, with $\hat{y}_t, \hat{\sigma}_t^2 = f^{\widehat{W}_t}(x)$, with T models and

$$\text{Var}(y) \approx \frac{1}{T} \sum_{i=1}^T \hat{y}_t^2 - \left(\frac{1}{T} \sum_{t=1}^T \hat{y}_t \right)^2 + \frac{1}{T} \sum_{t=1}^T \hat{\sigma}_t^2$$

- Objective : To minimise

$$\mathcal{L}_{BNN}(\theta) = \frac{1}{D} \sum_i \frac{1}{2} \hat{\sigma}_i^{-2} \|y_i - \hat{y}_i\|^2 + \frac{1}{2} \log \hat{\sigma}_i^2$$

Take home message

- In Bayesian deep learning, intractable distribution is approximated by KL divergence
- Dropout acts like approximate Bayesian inference in deep Gaussian processes.
- We can calculate model uncertainty without uncertain label
- We can combine aleatoric uncertainty with epistemic uncertainty

$$\mathcal{L}_{BNN}(\theta) = \frac{1}{D} \sum_i \frac{1}{2} \hat{\sigma}_i^{-2} \|y_i - \hat{y}_i\|^2 + \frac{1}{2} \log \hat{\sigma}_i^2$$